EVERYTOWN COMMUNITY SAFETY FUND
Organizational Capacity Index

OCI Capacity Coach

A 15-minute self-assessment that becomes a plan a CVI team can start Monday, and a live view of the whole network for CSF staff.

d2mi2h3eu8y69s.cloudfront.net · demo / demo · staff / staff

THE NEED

The 2025 pilot worked. Two people cannot hand-write reports for a whole network.

45

organizations took the manual Organizational Capacity Index pilot in 2025

2

CSF staff wrote every individualized report by hand

140+

grantee and alumni organizations waiting for the same support

Why it matters

Since 2019, CSF has granted $15.7 million to 141 community-based violence intervention programs in 72 cities. New grantees complete a self-assessment that shapes their training.

What the pilot taught us

Planning was the weakest area across the cohort. Planning, Fundraising and Financial Management move with overall capacity, so the Coach starts there.

Sources: Everytown OCI 2025 findings; everytownsupportfund.org

INTERVIEW WITH THE CSF TEAM

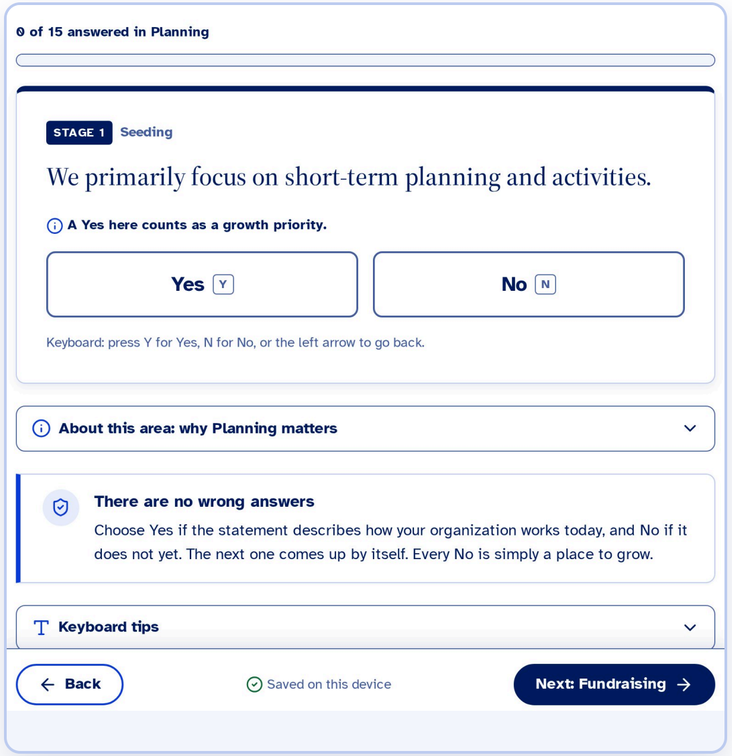
What Everytown told us, and what we built

THEY SAID	WE BUILT
“I can read this even though I'm at an 8th grade reading level and I've been to prison.”	**Plain language, measured.** Code checks the reading grade of every report and rewrites above grade 8.5. Read aloud, Spanish.
“Out of the 10 things you could fix, here's one or two priorities that you should do first.”	**Start here.** Code picks at most two first priorities. Then a 12-week roadmap, one week at a time.
“I know what training they need. I know what services they need.”	**Training needs.** Staff see which organizations need which training, and draft a training brief.
“Identify groups of grantees with similar problems that we could connect with each other.”	**Peers.** Similar grantees by program model, size and shared gaps, plus "Connect me" requests.

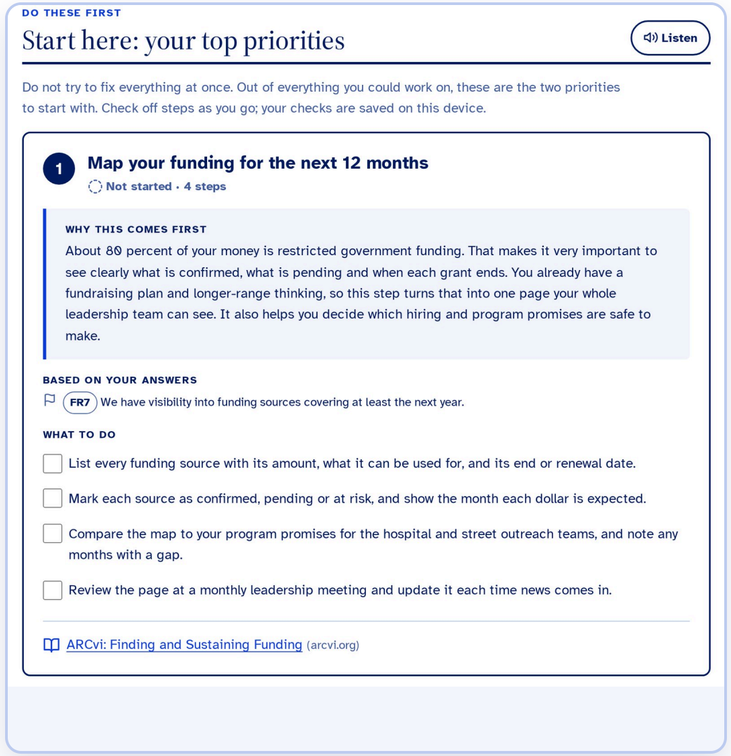
Quotes from our onsite interview with Everytown

LIVE DEMO · FOR THE ORGANIZATION'S LEADER

From Yes / No answers to a plan in about 40 seconds



1 **Answer.** One statement at a time, Yes or No, in English or Spanish. Teammates can answer an area.



2 **Start here.** Two priorities, the exact answers behind them and a vetted resource. Plus a PDF letter for the board.



3 **Follow through.** A 12-week roadmap with owners, points and badges. Week 12: retake the OCI.

LIVE DEMO · FOR CSF STAFF

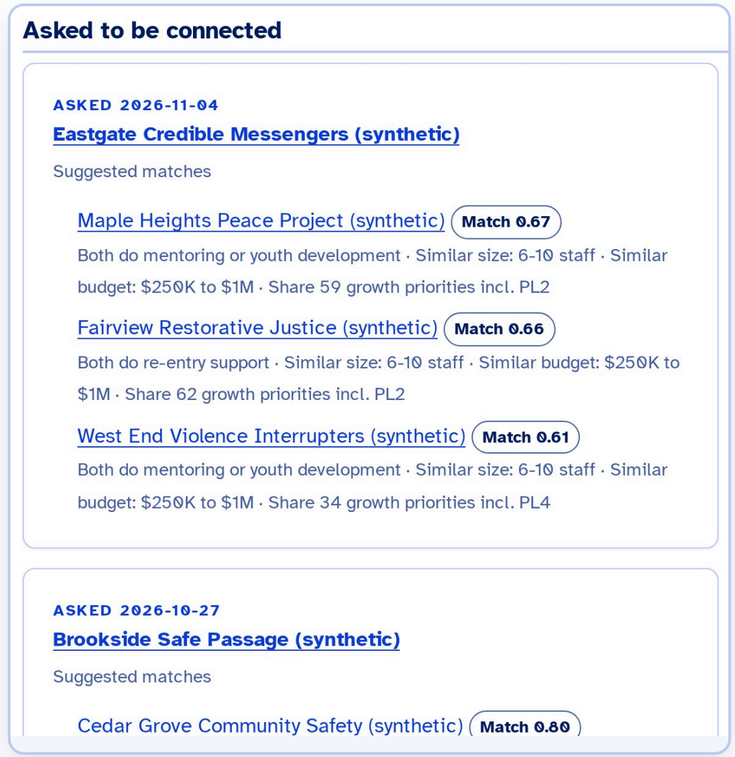
See the whole network, not one report at a time



Training needs. Who needs which training, who asked for help, and a draft training brief.



Ask the data. Plain questions. The AI picks a pre-built analysis; code computes every number.



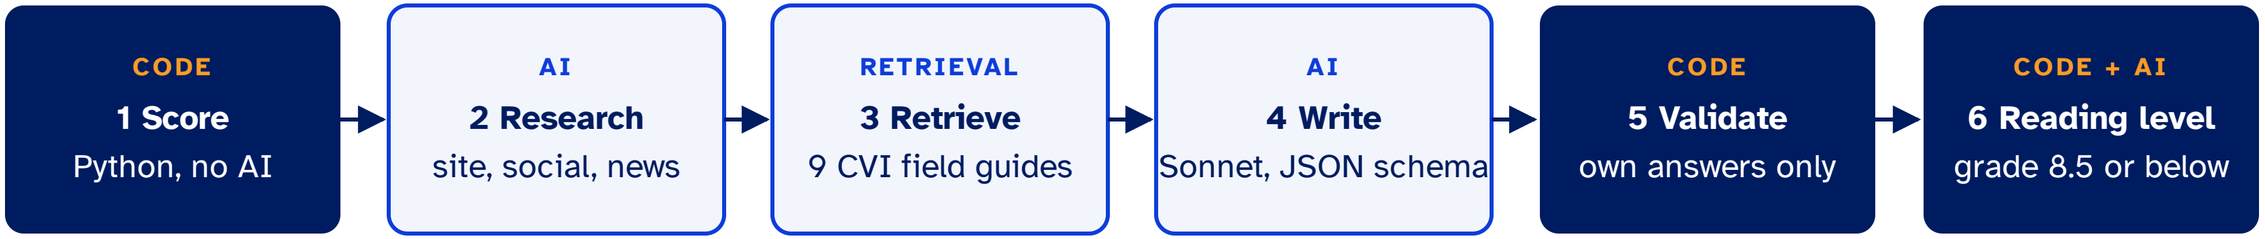
Peers. Similar grantees to connect, by program model, size and shared gaps.

Also: gaps by area, change over time ("Did support help?"), network report PDF, CSV exports

HOW IT WORKS

Code scores. AI writes. Code checks every line.

Six fixed steps, not an agent. One trace per report in Amazon CloudWatch.



Bedrock Guardrails screen untrusted text (websites, user questions). Crisis questions to the coach get a fixed 988 / 911 answer and never reach the model.

The AI never scores

The reference engine is unchanged and passes all 3 answer keys. Stages are final before any model runs.

Grounded or dropped

Every action must cite this organization's own answers. Anything else is removed before anyone sees it.

No invented links

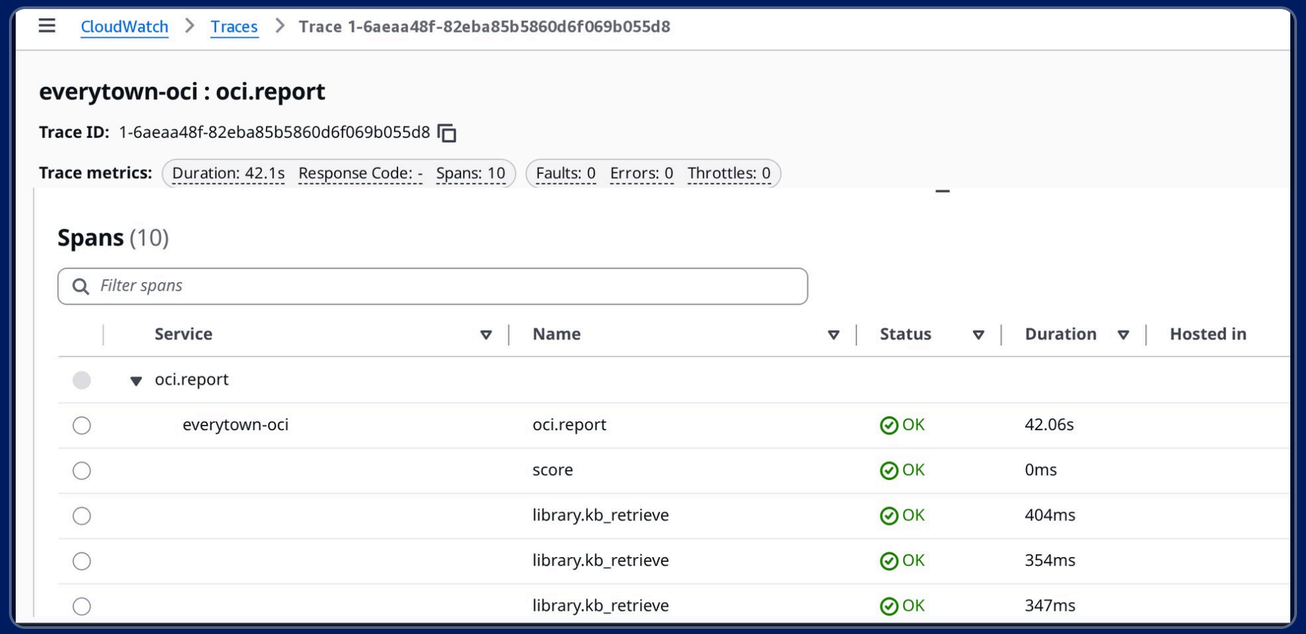
Resources only from a vetted 19-entry catalog of CVI guides.

Amazon CloudFront · **Amazon S3** · **AWS Lambda** · **Aurora PostgreSQL Serverless v2** · **Amazon Bedrock: Claude, Knowledge Bases, Guardrails** · **S3 Vectors** · **Amazon Polly** · **Amazon CloudWatch** · **AWS X-Ray** · **AWS CloudFormation**

Open source: github.com/ArielleTolome/everytown-oci

OBSERVABILITY AND EVALS

Every AI step is traced, tested and auditable



52 synthetic test cases

reports, coach questions and staff questions, scored by code checks and calibrated Claude judges. Every run lands in CloudWatch metrics and logs.

The evals caught 3 real bugs before any grantee did

- The guardrail blocked a normal Spanish report
- "Change my score" gave an error instead of a kind answer
- Tracing was silently off after a deploy

All in Amazon CloudWatch: Transaction Search traces, generative AI observability, Bedrock invocation logs, a dashboard and alarms. Staff open any report's trace from the dashboard.

Release gates: grounded, code-chosen focus, catalog links only, crisis and score-change handling, unsafe staff questions rejected

USER INTERFACE AND DESIGN

Written for the reader, not the funder

WCAG 2.2 AA
axe-core finds 0 violations on every page, at desktop and phone widths.

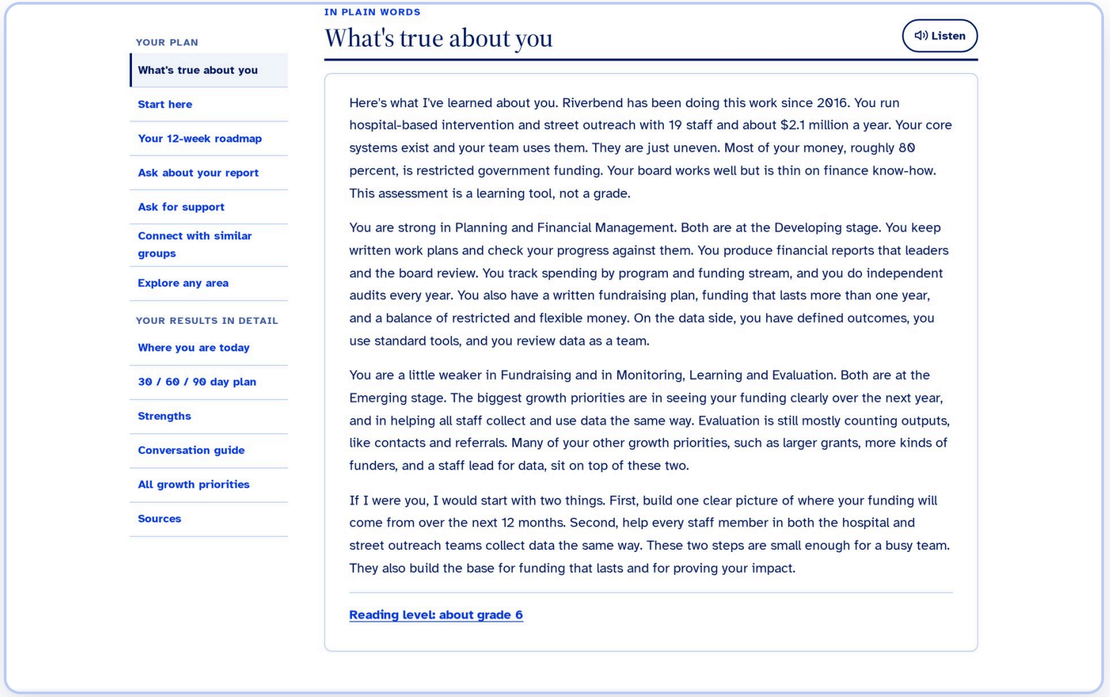
Grade 6 to 8
Reading level measured in code on every report, shown to the reader.

Strength first
A coach's letter: what is true about you, then growth priorities. Never a grade.

Listen
Any section read aloud with Amazon Polly. Larger text and high contrast in one menu.

English and Spanish
Questions, report and PDF, tagged so screen readers switch voice.

Everytown look
The CSF site's navy, serif headlines and pill buttons, at AA contrast.



Atkinson Hyperlegible, a typeface designed for low-vision readers

IMPACT AND SCALABILITY

The whole network, twice a year, for under $100 a year

40 s
from last answer to a finished AI report

$0.11
of Bedrock per report, from the trace's token counts

< $50
a year of tokens: 140 organizations, 2 assessments each

$1 to 3 / mo
hosting with Aurora auto-pause ($45 kept warm, as today)

Staff time goes to coaching

Reports write themselves. CSF staff spend the time on calls, trainings and convenings.

Progress over time

Week 12 is "Retake the OCI". Staff see stages move and whether support helped.

Forkable by Everytown

Open source, three CloudFormation stacks, one deploy script, serverless app.

Spend guards: 5 concurrent requests, an hourly report cap, optional AWS Budget. Observability: CloudWatch, a few dollars a month.

OUR ASK

Pilot with 3 to 4 CSF grantees this quarter

- CSF invites 3 to 4 grantees to take the assessment, about 15 minutes each.
- CSF staff review the reports with each organization and rate how helpful they were.
- Everytown Labs deploys the same templates in Everytown's own AWS account, with real sign-in.



Try it: d2mi2h3eu8y69s.cloudfront.net · demo / demo · staff / staff

Thank you to the Everytown Community Safety Fund team

github.com/ArielleTolome/everytown-oci

APPENDIX · RESPONSIBLE AI

Where the AI is allowed to act, and where it is not

Code decides

- ✓ Every score and stage (Everytown's reference engine)
- ✓ Which one or two priorities come first
- ✓ Which links appear (19-entry vetted catalog)
- ✓ Every number staff see in "Ask the data"
- ✓ Crisis questions: a fixed 988 / 911 answer
- ✓ Reading grade, and whether a rewrite is accepted

The AI writes, then gets checked

- ✓ The coach's letter and plan, through a JSON schema
- ✓ Short answers about your own report, cut at 200 words
- ✓ A 2 to 3 sentence summary of a table code built
- ✓ Output dropped if it cites answers you did not give
- ✓ Guardrails screen websites and user questions
- ✓ Blocked profile? The report goes on without it

Fixed steps, not agents: no model chooses tools, writes SQL or sees another organization's data. Judges in the evals only report; they never block or edit what users see.

APPENDIX · SECURITY AND PRIVACY

A sandbox build with production habits

Data

Synthetic organizations only. Aurora encrypted at rest, reached through the RDS Data API with an RDS-managed secret. Traces carry a role and a hashed session id, never a respondent's name.

Network

One CloudFront domain. The Lambda URL accepts only IAM-signed requests from that distribution. The website fetcher blocks private addresses. Traces and model logs stay in CloudWatch, behind AWS sign-in.

Spend and abuse

5 concurrent requests, an hourly cap on new reports, 15 coach questions per report, optional AWS Budget. Secrets stay out of git.

Before production

Real sign-in replaces the mocked demo logins (today anyone signed in can open a report by its id, acceptable only for synthetic data), CloudWatch Logs data protection to mask personal data in traces and model logs, and an Everytown retention policy for them.

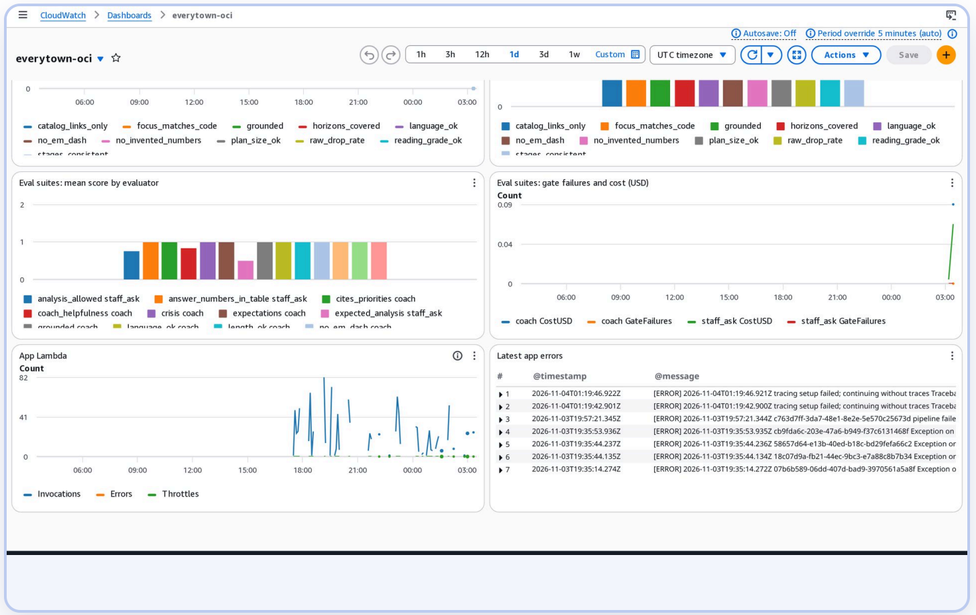
Teardown after judging: scripts/teardown.sh deletes every stack

APPENDIX · EVALS

What we test on every release

v1.1 baseline; every run publishes to CloudWatch. Judge scores from 0 to 1.

Suite	Cases	Gate	Judges
Reports EN/ES samples, edge cases, prompt injection	16	Pass	personalization 0.88, actionability 0.84, faithfulness 0.75, plain language 0.72
Coach how-to, change my score, crisis, off-topic	21	Pass	helpfulness 0.69
Staff questions including SQL injection and delete requests	15	Pass	all unsafe requests rejected
Judge calibration good vs deliberately bad outputs	20	Pass	every judge ranks bad below good, twice in a row



python3 -m evals.run --suite all · --compare runs a model A/B · exits nonzero when a gate fails

APPENDIX · COST

What it costs to run

Item	Cost	Notes
One AI report	about $0.11	Claude Sonnet with a cached prompt, plus Haiku research; measured per trace
Whole network, twice a year	under $50 a year	140 organizations × 2 assessments, with room for retries and coach questions
App hosting	$1 to $3 a month	Aurora on auto-pause; about $45 a month kept warm at 0.5 ACU (as for judging)
CloudWatch observability	a few dollars a month	dashboard, a few custom metrics, spans and model logs at demo volume; no server
A full eval run	about $3.50	all three suites with judges; a quick check is about $0.60

Costs from the live deployment and CloudWatch traces, us-west-2

APPENDIX · ROADMAP

Built to be iterated on, not rebuilt

Live today

- ✓ All 8 areas, 143 statements
- ✓ Report, PDF letter, 12-week roadmap
- ✓ Staff dashboard, Ask the data, Peers
- ✓ English and Spanish
- ✓ Tracing and evals in CloudWatch

Next, with Everytown

- ✓ Pilot with 3 to 4 grantees
- ✓ Real sign-in in the Grantee Portal
- ✓ CSF experts review the field guides
- ✓ Add pilot feedback to the eval sets

Later, as data grows

- ✓ Peer benchmarking beyond the staff preview
- ✓ Match new grantees with established peers
- ✓ More languages